

Accelerated DataScience Tools Overview

On choosing the right tool between **Pandas**, Modin, Dask, CuDF, SQLite, Spark, **NumPy**, CuPy, **NetworkX**, CuGraph and RetworkX

--

By Ashot Vardanian

Founder @ [Unum.Cloud](#)

linkedin.com/in/ashvardanian

fb.com/ashvardanian

t.me/ashvardanian

Hardware

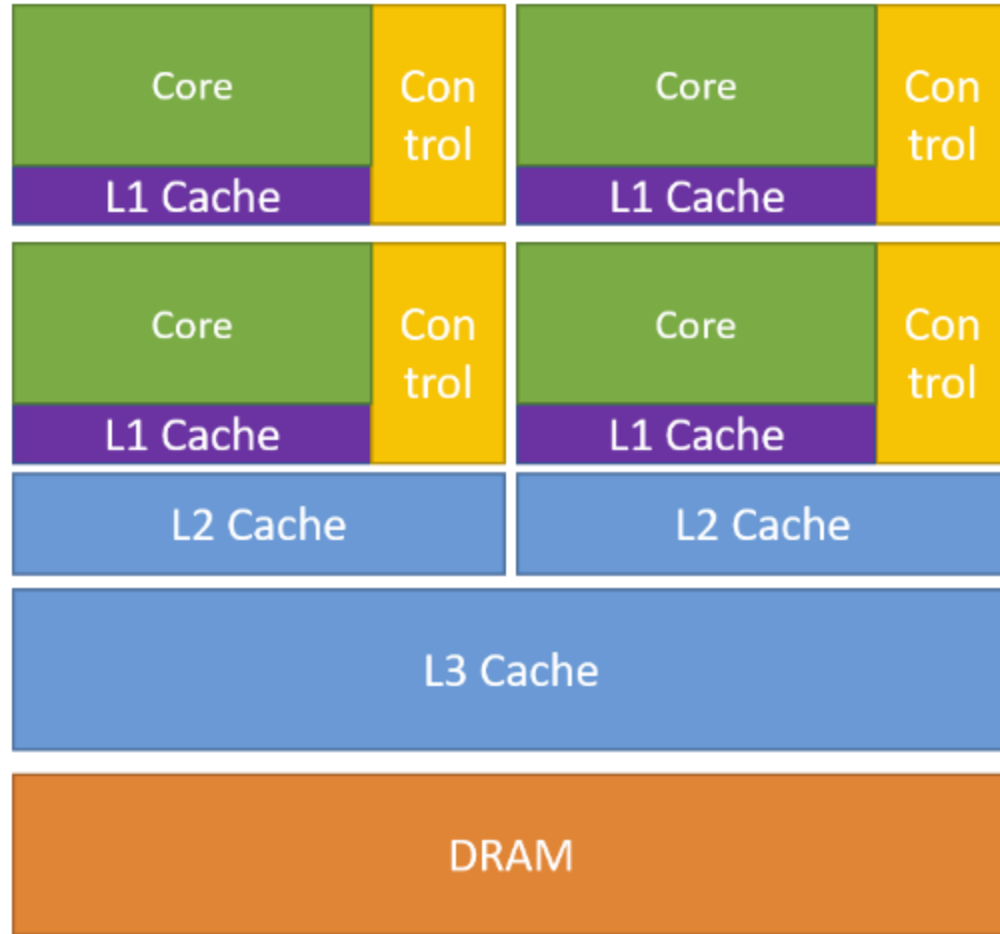
- Multi-core CPUs
- Highly-parallel GPUs

100 - 10'000 threads/device.


200+ Gbit networking.

Datasets 100x bigger than RAM.

CPU vs GPU



CPU



GPU

A100 Zoom In



How to sum some numbers?

Python:

```
sum(x)
```

C++:

```
std::accumulate(x.begin(), x.end(), 0.f);
```

Accumulation, the hardest task!


```
__global__ void reduce_warps(float const *inputs, unsigned int input_size, float *outputs) {
    float sum = 0;
    for (unsigned int i = blockIdx.x * blockDim.x + threadIdx.x; i < input_size; i += blockDim.x * gridDim.x)
        sum += inputs[i];

    __shared__ float shared[32];
    unsigned int lane = threadIdx.x % warpSize;
    unsigned int wid = threadIdx.x / warpSize;


    sum = reduce_warp(sum); // Important
    if (lane == 0)
        shared[wid] = sum;
    __syncthreads();

    sum = (threadIdx.x < blockDim.x / warpSize) ? shared[lane] : 0;
    if (wid == 0)
        sum = reduce_warp(sum); // Important
    if (threadIdx.x == 0)
        outputs[blockIdx.x] = sum;
}
```

What's inside?



What's the alternative?



Why Suffer?

Python: **1 GB/s**

```
sum(x)
```

C++ 17: **87 GB/s**

```
std::reduce(std::execution::par_unseq, x.begin(), x.end(), 0.f);
```

CUDA: **789 GB/s**

Can Performance Be Usable?

Python:

```
sum(x)
```

C++:

```
std::accumulate(x.begin(), x.end(), 0.f);
```

CUDA + Thrust:

```
thrust::reduce(x.begin(), x.end(), 0.f);
```

If C++ can be readable, can Python be FAST?

- NumPy → CuPy
- NetworkX → CuGraph or RetworkX
- Pandas →
 - Modin is Multi-Threaded
 - Dask is Multi-Node
 - CuDF is on GPUs
 - Dask-CuDF is Multi-Node on GPUs

Change is Hard Easy

```
# import numpy as np
import cupy as np

np.matmul(mat, mat)
```

Yields us almost compatible API!

```
# np.random.rand(100, 100).astype(np.float32)
cupy.random.rand(100, 100, dtype=np.float32)
cupy.cuda.stream.get_current_stream().synchronize()
```

Calm your Horses! CPUs have something to say...

Let's check our configs:

```
name: benchmark
channels:
  - conda-forge
  - defaults
dependencies:
  - numpy
```

Do you know what you are getting?

The Precise Way

```
dependencies:  
- 'blas*=mkl'  
- numpy
```

or:

```
dependencies:  
- 'blas*=openblas'  
- numpy
```

What Have We Achieved?

In short, up to 1000x performance improvements!

| Speedups | 512 ² | 1024 ² | 2048 ² | 4096 ² | 8192 ² | 16384 ² |
|----------------|------------------|-------------------|-------------------|-------------------|-------------------|--------------------|
| SVD | 0.8x | 0.7x | 0.6x | 0.5x | 0.5x | 0.6x |
| Pearson Corr. | 2.5x | 2.0x | 1.4x | 1.1x | 1.5x | 1.4x |
| GEMM | 8.3x | 17.9x | 25.1x | 21.0x | 19.9x | 23.4x |
| Moving Average | 10.9x | 19.9x | 49.0x | 59.3x | 59.4x | 53.9x |
| Medians | 11.2x | 30.3x | 58.8x | 72.2x | 99.4x | 83.5x |
| Sorting | 87.7x | 274.4x | 547.3x | 788.8x | 920.5x | 1008.5x |

Tell us more! What's up with Graphs?

- NetworkX = Python
- RetworkX = Rust 🤢
- CuGraph = CUDA 🔥

How close are they?

```
networkx.weakly_connected_components(g)
```

```
retworkx.weakly_connected_components(g)
```

```
cugraph.components.connectivity.weakly_connected_components(g)
```

How close are they? Reality

```
networkx.pagerank(g)
```

```
raise NotImplemented()
```

```
cugraph.pagerank(g)
```


How close are they? Algorithm Mismatch

```
networkx.spring_layout(g)
```

```
retworkx.spring_layout(g)
```


```
cugraph.force_atlas2(g)
```

What Are We Fighting For? Speed!




For Bitcoin Graph: **124x** improvement!

PageRank



On big graphs: **247x** improvement!

No Guarantees



Stack: 1.6 GB.

Orkut: 1.7 GB.

What about Tabular Data?

Let's take the [Taxi Rides Dataset](#)!

```
aws s3 ls --recursive s3://ursa-labs-taxi-data/ --recursive --human-readable --summarize  
aws s3 sync s3://ursa-labs-taxi-data/ ADSB
```

Or in R:

```
arrow::copy_files("s3://ursa-labs-taxi-data", "nyc-taxi")
```

Producing 40 GB in clean Parquet files.

What will we do? SQL time!

Query 1 in SQL:

```
SELECT cab_type,  
       count(*)  
FROM trips  
GROUP BY 1;
```

In Pandas:

```
selected_df = trips[['cab_type']]  
grouped_df = selected_df.groupby('cab_type')  
final_df = grouped_df.size().reset_index(name='counts')
```


Query 2: Average by Group

```
SELECT passenger_count,  
       avg(total_amount)  
FROM trips  
GROUP BY 1;
```

In Pandas:

```
selected_df = trips[['passenger_count', 'total_amount']]  
grouped_df = selected_df.groupby('passenger_count')  
final_df = grouped_df.mean().reset_index()
```

Query 3: Transform & Histogram

```
SELECT passenger_count,  
       extract(year from pickup_datetime),  
       count(*)  
FROM trips  
GROUP BY 1,  
        2;
```

Our dataset contains dates in the following format: "2020-01-01 00:35:39".

```
selected_df = trips[['passenger_count', 'pickup_datetime']]  
selected_df['year'] = pd.to_datetime(  
    selected_df.pop('pickup_datetime'),  
    format='%Y-%m-%d %H:%M:%S'  
) .dt.year  
grouped_df = selected_df.groupby(['passenger_count', 'year'])  
final_df = grouped_df.size().reset_index(name='counts')
```

Query 4: All Together

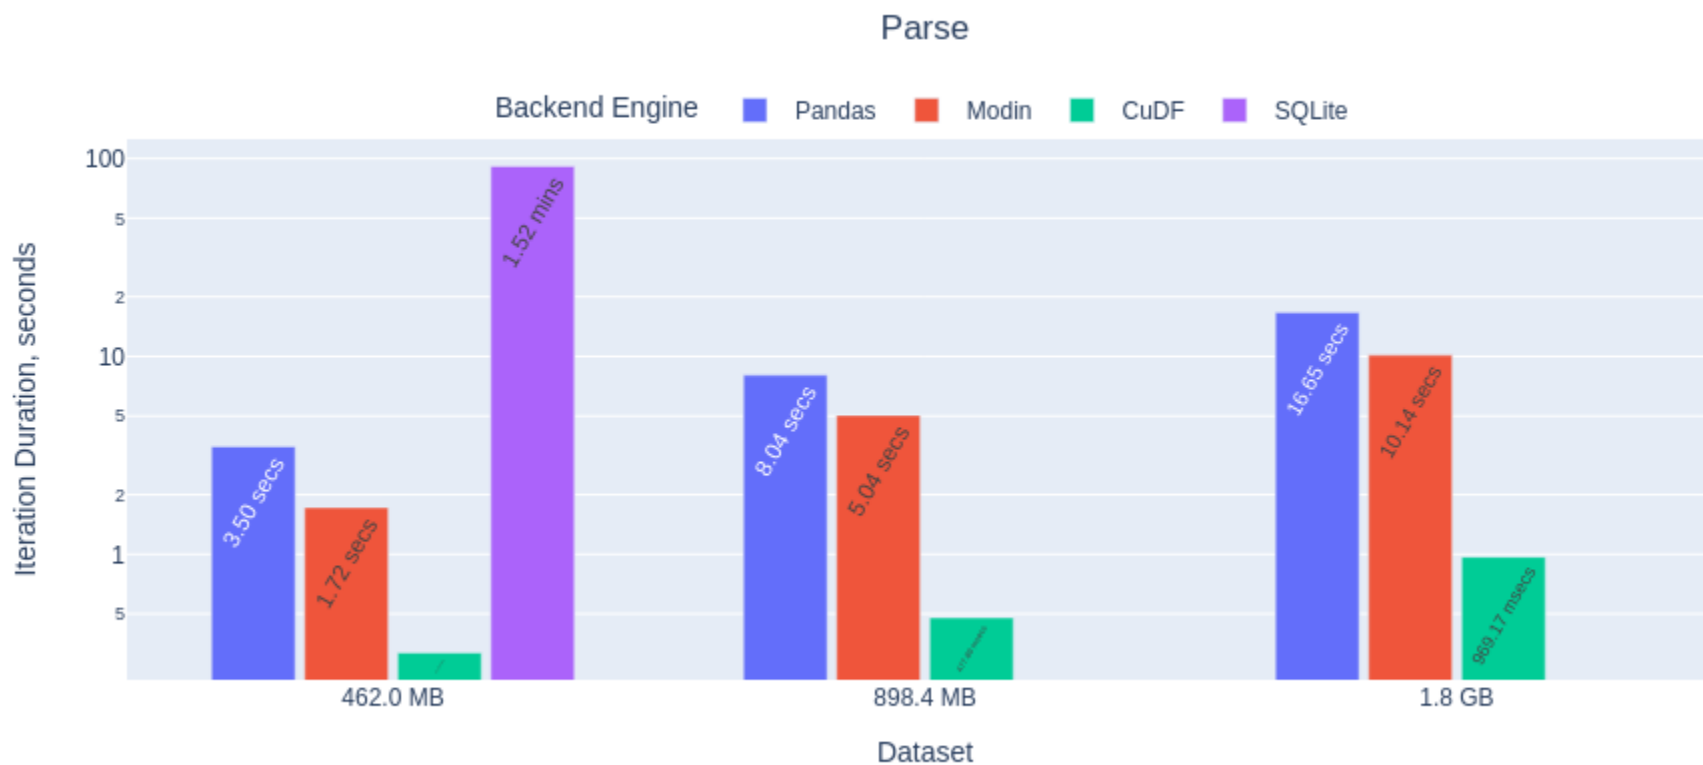
```
SELECT passenger_count,  
       extract(year from pickup_datetime),  
       round(trip_distance),  
       count(*)  
FROM trips  
GROUP BY 1,  
        2,  
        3  
ORDER BY 2,  
        4 desc;
```

```
selected_df = trips[['passenger_count', 'pickup_datetime', 'trip_distance']]  
selected_df['trip_distance'] = selected_df['trip_distance'].round().astype(int)  
selected_df['year'] = pd.to_datetime(selected_df.pop('pickup_datetime'), format='%Y-%m-%d %H:%M:%S').dt.year  
grouped_df = selected_df.groupby(['passenger_count', 'year', 'trip_distance'])  
final_df = grouped_df.size().reset_index(name='counts')  
final_df = final_df.sort_values(['year', 'counts'], ascending=[True, False])
```

Porting to Pandas and beyond!

- Pandas supports `reset_index(name='')` on series, but not on frames. Other libraries mostly don't have that so we rename afterwards for higher compatibility.
- In queries 3 and 4 we could have fetched/converted data from the main source in just a single run, but to allow lazy evaluation of `WHERE`-like sampling queries, we split it into two steps.
- Major problem in Dask is the lack of compatible constructors, the most essential function of any class. You are generally expected to start with Pandas and cuDF and later [convert those](#).

Results: Parsing



Pandas compatiability: CuDF , Modin 

Results: Query 1



Pandas compatiability: CuDF ✓, Modin ✓

Results: Query 2



Pandas compatiability: CuDF ✓, Modin ✓

Results: Query 3



Pandas compatiability: CuDF ✓, Modin ✓

Results: Query 4



Pandas compatiability: CuDF ✓, Modin ✓

What's next?

- Extending NumPy comparisons: + JAX, + ATen.
- Publishing more tabular results: Dask, Spark.
- Machine Learning benchmarks: PT, TF, JAX.

General Recommendations

- Always provide `dtype` , don't use `float` .
- Power-of-two array sizes.
- Prefer symmetrical tensors.
- Structure-of-Arrays instead of Array-of-Structures.

Known Pitfalls

- Multi-GPU.
- Unified Memory Latencies.
- NV-Link vs NUMA nodes and AMD MI200.
- Synchronization of CUDA streams and graphs.

Where we apply?

- Processing 50 GB/s on a single machine.
- Rapid experimentation with datasets between 100 MB and 4 GB.
- Fast experimentation with dataset above 4 GB.

All at [Unum.cloud](https://unum.cloud)!

--

linkedin.com/in/ashvardanian

fb.com/ashvardanian

t.me/ashvardanian